

# CSCD 439/539 GPU Computing HW4

## Modify Reduction

No Late Submissions are accepted. **Rules:** Your code must CUDA C Language. If your program shows a compilation or run-time error, you get a zero for this assignment.

**Submission:** Wrap up all your **source files** into a single zip file. Name your zip file as *FirstInitialYourLastNameHw4.zip*. For example, if your legal name is Will Smith, you should name your zip file as wsmithHw4.zip. A simple makefile has been provided in the zip file.

**For archive purpose, please also submit your single zip file on EWU Canvas by following CSCD439-01 Course Assignments Hw4 Submit Assignment to upload your single zip file.**

### Problem Description:

Based on the lecture about reduction algorithm on CUDA device, you are required to implement the following features and answer the questions.

1, Download the demo code d10\_reduce4\_add.tar.gz on canvas under Files DemoCode. Read and understand the provided CUDA C program. We have already gone through this code in class. The demo code is also included in the start up package.

2, The provided kernel in the demo code, `__global__ void reduce2(float *in, float *out, int n)`, is not optimized. The basic idea about kernel `reduce2()` is described in the lecture notes regarding reduction algorithm. The following diagram (figure 1) illustrates `reduce2()` kernel.

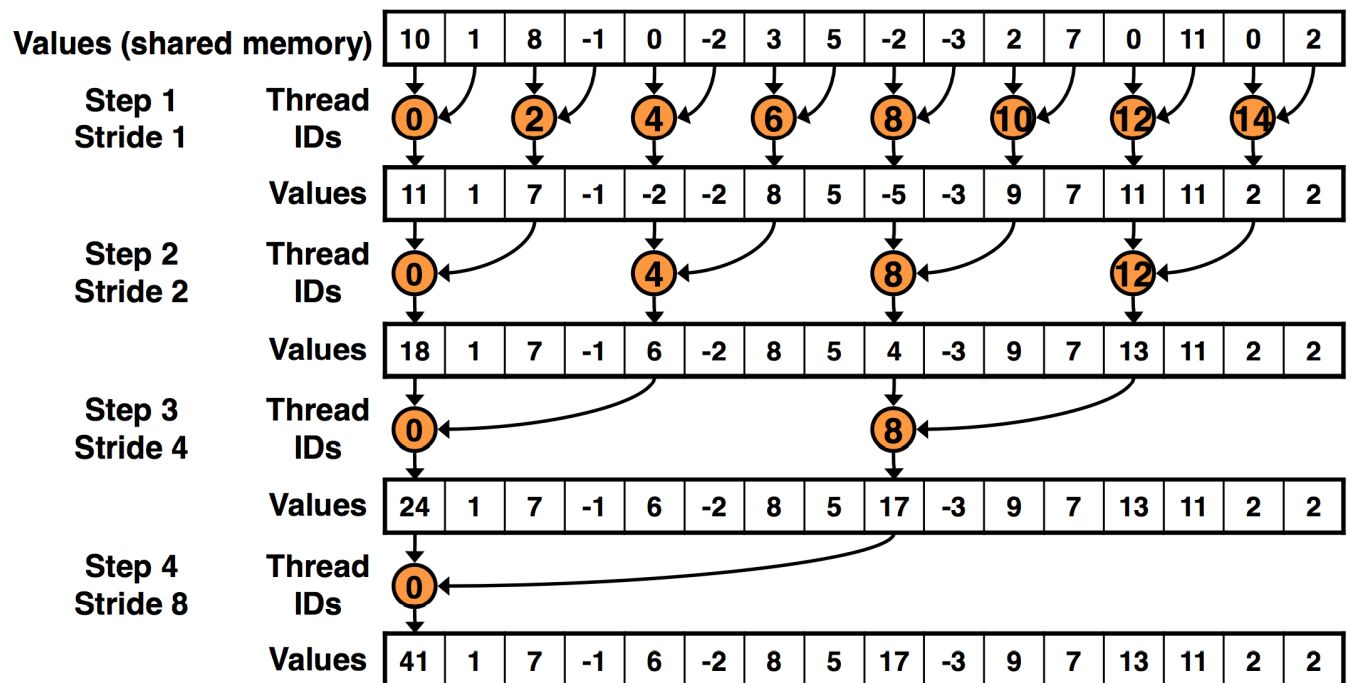


Figure 1

3, You have to write another kernel `reduce3()`, that also performs reduction on CUDA device. Kernel `reduce3()` accepts the same set of parameters list as `reduce2()` kernel does, with each parameter maintaining the same meaning. But `reduce3()` kernel uses the following threads-data mapping and data access patterns, as described below in figure 2.

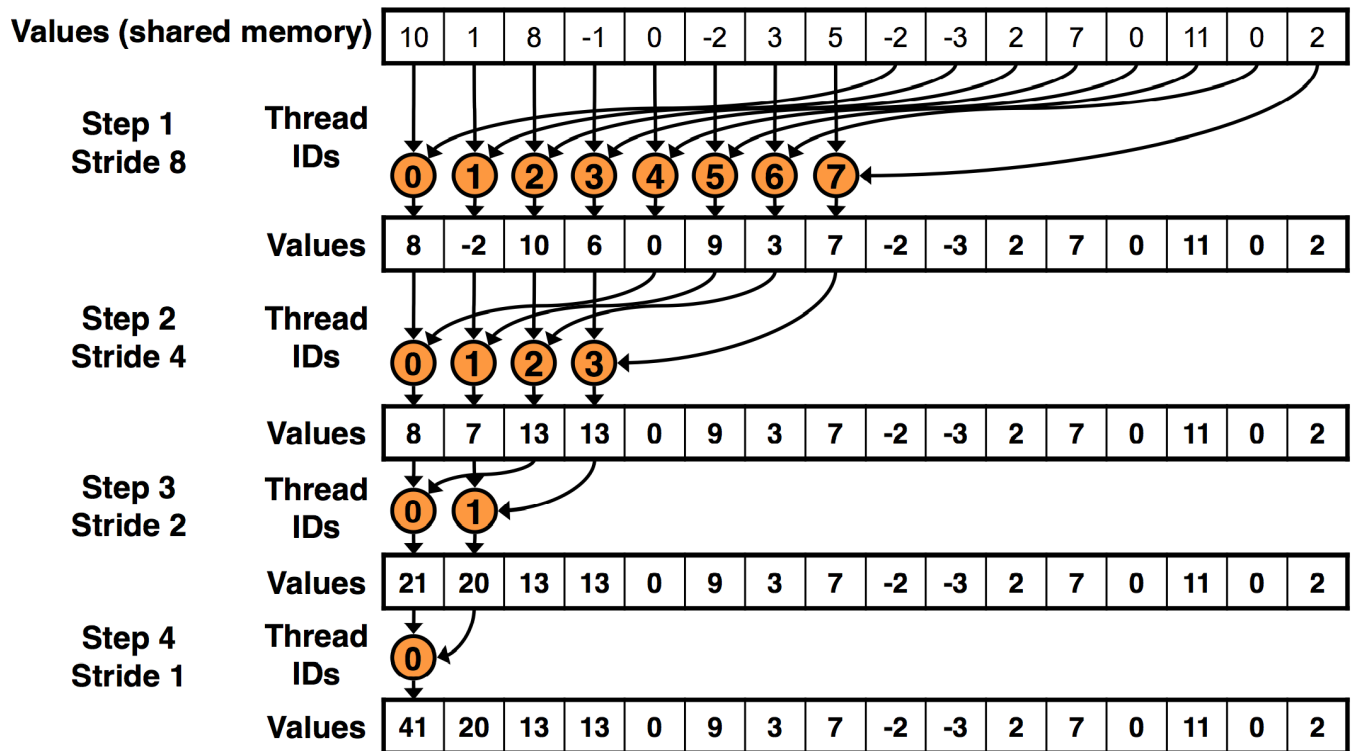


Figure 2

In the design of figure 2, **the size of thread block (in a shape of 1 \* N) is required** to be half of the tile size, (i.e. also half of the thread block size used in the kernel `reduce2()`). In comparison, `reduce3()` will use less number of threads in each block, but can achieve the same goal as `reduce2()`. In this way, we improve the performance of `reduce3()`, compared with `reduce2()`. **You have to change the execution configuration parameters in the main()** in order to set up the new block size variable for `reduce3()`. Another major difference between `reduce2()` and `reduce3()` kernel pertains to which pair of data elements are processed by a thread.

3, Change the `main()` function to invoke the kernel **`reduce3()`** you implemented, and measure the kernel execution time. Then you have to compare the kernel `reduce3()` with the `reduce2()` by filling out the following table.

Input size	1048576	16777216	67108864	134217728
Block Dimensions	1 X 1024	1 X 1024	1 X 1024	1 X 1024
T1:time cost for reduce2 (ms)				
T2:time cost for reduce3 (ms)				
Speedup = T1 / T2				

4, Do you identify any advantage of reduce3() over reduce2() kernel? Or vice versus? And why ? ( hint: in terms of performance, such as bank conflicts, condition divergence or idle threads. )